

# homework-07

August 4, 2025

## 1 Homework 7

### 1.1 References

- Lectures 24-26 (inclusive).

### 1.2 Instructions

- Type your name and email in the “Student details” section below.
- Develop the code and generate the figures you need to solve the problems using this notebook.
- For the answers that require a mathematical proof or derivation you should type them using latex. If you have never written latex before and you find it exceedingly difficult, we will likely accept handwritten solutions.
- The total homework points are 100. Please note that the problems are not weighed equally.

```
[ ]: import matplotlib.pyplot as plt
%matplotlib inline
import matplotlib_inline
matplotlib_inline.backend_inline.set_matplotlib_formats('svg')
import seaborn as sns
sns.set_context("paper")
sns.set_style("ticks")

import scipy
import numpy as np
import scipy.stats as st
import urllib.request
import os

def download(
    url : str,
    local_filename : str = None
):
    """Download a file from a url.

    Arguments
    url           -- The url we want to download.
    local_filename -- The filename to write on. If not
                    specified
```

```

"""
if local_filename is None:
    local_filename = os.path.basename(url)
urllib.request.urlretrieve(url, local_filename)

```

### 1.3 Student details

- **First Name:** Hannah
- **Last Name:** Moskios
- **Email:** hmoskios@purdue.edu

In this problem, you must use a deep neural network (DNN) to perform a regression task. The dataset we are going to use is the [Airfoil Self-Noise Data Set](#). From this reference, the description of the dataset is as follows:

The NASA data set comprises different size NACA 0012 airfoils at various wind tunnel speeds and angles of attack. The span of the airfoil and the observer position were the same in all of the experiments.

**Attribute Information:** This problem has the following inputs: 1. Frequency, in Hertz. 2. The angle of attack, in degrees. 3. Chord length, in meters. 4. Free-stream velocity, in meters per second. 5. Suction side displacement thickness, in meters.

The only output is: 6. Scaled sound pressure level in decibels.

You will have to do regression between the inputs and the output using a DNN. Before we start, let's download and load the data.

```
[ ]: !curl -O --insecure "https://archive.ics.uci.edu/ml/machine-learning-databases/
↪00291/airfoil_self_noise.dat"
```

% Total	% Received	% Xferd	Average Speed	Time	Time	Time	Current
			Dload Upload	Total	Spent	Left	Speed
0 0 0 0 0 0 0 0	--:--:-- --:--:-- --:--:-- 0						
100 59984 0 59984 0 0 294k 0	--:--:-- --:--:-- --:--:-- 295k						

The data are in simple text format. Here is how we can load them:

```
[ ]: data = np.loadtxt('airfoil_self_noise.dat')
data
```

```
[ ]: array([[8.00000e+02, 0.00000e+00, 3.04800e-01, 7.13000e+01, 2.66337e-03,
   1.26201e+02],
 [1.00000e+03, 0.00000e+00, 3.04800e-01, 7.13000e+01, 2.66337e-03,
   1.25201e+02],
 [1.25000e+03, 0.00000e+00, 3.04800e-01, 7.13000e+01, 2.66337e-03,
   1.25951e+02],
 ...,
 [4.00000e+03, 1.56000e+01, 1.01600e-01, 3.96000e+01, 5.28487e-02,
   1.06604e+02],
 [5.00000e+03, 1.56000e+01, 1.01600e-01, 3.96000e+01, 5.28487e-02,
```

```

1.06224e+02],
[6.30000e+03, 1.56000e+01, 1.01600e-01, 3.96000e+01, 5.28487e-02,
1.04204e+02]])

```

You may work directly with `data`, but, for your convenience, I am going to put them also in a nice Pandas DataFrame:

```
[ ]: import pandas as pd
df = pd.DataFrame(data, columns=['Frequency', 'Angle_of_attack', 'Chord_length',
                                  'Velocity', 'Suction_thickness', ↴
                                  'Sound_pressure'])
df
```

	Frequency	Angle_of_attack	Chord_length	Velocity	Suction_thickness	Sound_pressure
0	800.0	0.0	0.3048	71.3	0.002663	126.201
1	1000.0	0.0	0.3048	71.3	0.002663	125.201
2	1250.0	0.0	0.3048	71.3	0.002663	125.951
3	1600.0	0.0	0.3048	71.3	0.002663	127.591
4	2000.0	0.0	0.3048	71.3	0.002663	127.461
...	...	...	...	...	...	...
1498	2500.0	15.6	0.1016	39.6	0.052849	110.264
1499	3150.0	15.6	0.1016	39.6	0.052849	109.254
1500	4000.0	15.6	0.1016	39.6	0.052849	106.604
1501	5000.0	15.6	0.1016	39.6	0.052849	106.224
1502	6300.0	15.6	0.1016	39.6	0.052849	104.204
[1503 rows x 6 columns]						

### 1.3.1 Part A - Analyze the data visually

It is always a good idea to visualize the data before you start doing anything with them.

**Part A.I. - Do the histograms of all variables** Use as many code segments as you need below to plot the histogram of each variable (all inputs and the output in separate plots) Discuss whether or not you need to standardize the data before moving to regression.


**Answer:**

```
[ ]: # Your code here

# Create a histogram of all the variables
axes = df.hist(bins=30, figsize=(10, 10));
axes = axes.ravel()

# Add labels to the x-axis and y-axis
columns=['Frequency (Hz)', 'Angle of Attack (Degrees)', 'Chord Length (m)',
          'Velocity (m/s)', 'Suction Thickness (m)', 'Sound Pressure (dB)']
for ax, col in zip(axes, columns):
    ax.set_xlabel(col)
    ax.set_ylabel('Count')

# Formatting
plt.tight_layout()
sns.despine(trim=True)
```



The histograms show that each variable occupies a different numerical scale. For example, suction thickness ranges from 0-0.06, whereas sound pressure ranges from 100-140. Therefore, we must standardize the data before moving on to regression.

**Part A.II - Do the scatter plots between all input variables** Do the scatter plot between all input variables. This will give you an idea of the range of experimental conditions. Whatever model you build will only be valid inside the domain implicitly defined with your experimental conditions. Are there any holes in the dataset, i.e., places where you have no data?

**Answer:**

[7]: # Your code here

# Number of inputs and outputs

```


n_outputs = 1
n_inputs = df.shape[1] - 1

# Initialize the subplots
fig, ax = plt.subplots(n_inputs, n_inputs, figsize=(10, 10), dpi=300);

# Plot the scatter plots between all input variables
for i in range(n_inputs):
    for j in range(n_inputs):
        ax[i, j].scatter(df.iloc[:, j], df.iloc[:, i], alpha=0.5,
                          edgecolor = 'k', rasterized=True)
        ax[i, j].grid(True)
        if j == 0:
            ax[i, j].set_ylabel(df.columns[i])
        if i == n_inputs - 1:
            ax[i, j].set_xlabel(df.columns[j])
        if i != n_inputs - 1:
            ax[i, j].xaxis.set_ticklabels([])
        if j != 0:
            ax[i, j].yaxis.set_ticklabels([])

# Display the scatter plots
plt.show()

```



There are some places where we do not have any data. For example, velocity and chord length have large gaps between each cluster of data points. This is because there are only 4 unique velocity values and 6 unique chord length values in the data. In contrast, the other input variables have fewer holes in the data. Although the suction thickness, angle of attack, and frequency have a few missing values, there is data for the majority of the range.

**Part A.III - Do the scatter plots between each input and the output** Do the scatter plot between each input variable and the output. This will give you an idea of the functional relationship between the two. Do you observe any obvious patterns?


**Answer:**

```
[ ]: # Your code here
plt.figure(figsize=(10, 7))
for i in range(n_inputs):
```

```

plt.subplot(2, 3, i + 1)
plt.scatter(df.iloc[:, i], df.iloc[:, -1], alpha=0.5, edgecolor = 'k')
plt.xlabel(columns[i])
plt.ylabel(columns[-1])
plt.tight_layout()

```



There are no obvious patterns between the input and output variables. On the scatter plots, the data points appear diffuse and there is no clear trend. As the frequency and suction thickness increase, it is possible that the sound pressure decreases slightly. However, this is not obvious and the correlation is very weak.

As in part A.II, we see the large gaps in the data for velocity and chord length.

### 1.3.2 Part B - Use DNN to do regression

Let start by separating inputs and outputs for you:

```
[ ]: X = data[:, :-1]
y = data[:, -1][:, None]
```

**Part B.I - Make the loss** Use standard torch functionality to create a function that gives you the sum of square error followed by an L2 regularization term for the weights and biases of all

network parameters (remember that the L2 regularization is like putting a Gaussian prior on each parameter). Follow the instructions below and fill in the missing code.

**Answer:**

```
[ ]: import torch
import torch.nn as nn

# Use standard torch functionality to define a function
# mse_loss(y_obs, y_pred) which gives you the mean of the sum of the square
# of the difference between y_obs and y_pred
# Hint: This is already implemented in PyTorch. You can just reuse it.
mse_loss = nn.MSELoss() # your code here
```

```
[ ]: # Test your code here
y_obs_tmp = np.random.randn(100, 1)
y_pred_tmp = np.random.randn(100, 1)
print('Your mse_loss: {0:1.2f}'.format(mse_loss(torch.Tensor(y_obs_tmp),
                                                torch.Tensor(y_pred_tmp))))
print('What you should be getting: {0:1.2f}'.format(np.mean((y_obs_tmp -
    y_pred_tmp) ** 2)))
```

Your mse\_loss: 1.50  
What you should be getting: 1.50

```
[ ]: # Now, we will create a regularization term for the loss
# I'm just going to give you this one:
def l2_reg_loss(params):
    """
    This needs an iterable object of network parameters.
    You can get it by doing `net.parameters()`.

    Returns the sum of the squared norms of all parameters.
    """
    l2_reg = torch.tensor(0.)
    for p in params:
        l2_reg += torch.norm(p) ** 2
    return l2_reg
```

```
[ ]: # Finally, let's add the two together to make a mean square error loss
# plus some weight (which we will call reg_weight) times the sum of the squared
# norms of all parameters.
# I give you the signature and you have to implement the rest of the code:
def loss_func(y_obs, y_pred, reg_weight, params):
    """
    Parameters:
    y_obs      -      The observed outputs
    y_pred     -      The predicted outputs
    """
```

```

reg_weight -      The regularization weight (a positive scalar)
params       -      An iterable containing the parameters of the network

Returns the sum of the MSE loss plus reg_weight times the sum of the
squared norms of all parameters.

"""
# Your code here
# MSE loss
MSE_loss = mse_loss(y_obs, y_pred)

# L2 regularization term
L2_reg = l2_reg_loss(params)

# Total loss
total_loss = MSE_loss + reg_weight * L2_reg

return total_loss

```

```

[ ]: # You can try your final code here
# First, here is a dummy model
dummy_net = nn.Sequential(nn.Linear(10, 20),
                          nn.Sigmoid(),
                          nn.Linear(20, 1))
loss = loss_func(torch.Tensor(y_obs_tmp), torch.Tensor(y_pred_tmp),
                 0.0,
                 dummy_net.parameters())
print('The loss without regularization: {:.2f}'.format(loss.item()))
print('This should be the same as this: {:.2f}'.format(mse_loss(torch.
    Tensor(y_obs_tmp), torch.Tensor(y_pred_tmp))))
loss = loss_func(torch.Tensor(y_obs_tmp), torch.Tensor(y_pred_tmp),
                 0.01,
                 dummy_net.parameters())
print('The loss with regularization: {:.2f}'.format(loss.item()))

```

The loss without regularization: 1.50

This should be the same as this: 1.50

The loss with regularization: 1.58

**Part B.III - Write flexible code to perform regression** When training neural networks, you must hand-pick many parameters, from the network structure to the activation functions to the regularization parameters to the details of the stochastic optimization. Instead of mindlessly going through trial and error, it is better to think about the parameters you want to investigate (vary) and write code that allows you to train networks with all different parameter variations repeatedly. In what follows, I will guide you through writing code for training an arbitrary regression network having the flexibility to:

- standardize the inputs and output or not
- experiment with various levels of regularization

- change the learning rate of the stochastic optimization algorithm
- change the batch size of the optimization algorithm
- change the number of epochs (how many times the optimization algorithm does a complete sweep through all the data).

**Answer:**

```
[ ]: # We will start by creating a class that encapsulates a regression
# network so that we can turn on or off input/output standardization
# without too much fuss.
# The class will represent a trained network model.
# It will "know" whether or not during training we standardized the data.
# I am not asking you to do anything here, so you can run this code segment
# or read through it if you want to know the details.
from sklearn.preprocessing import StandardScaler

class TrainedModel(object):
    """
    A class that represents a trained network model.
    The main reason I created this class is to encapsulate the standardization
    process in an excellent way.

    Parameters:

    net          -   A network.
    standardized -   True if the network expects standardized features and
                     outputs standardized targets. False otherwise.
    feature_scaler - A feature scalar - Ala scikit.learn. Must have
                      transform() and inverse_transform() implemented.
    target_scaler - Similar to feature_scaler but for targets...
    """

    def __init__(self, net, standardized=False, feature_scaler=None, ↴
                 target_scaler=None):
        self.net = net
        self.standardized = standardized
        self.feature_scaler = feature_scaler
        self.target_scaler = target_scaler

    def __call__(self, X):
        """
        Evaluates the model at X.
        """
        # If not scaled, then the model is just net(X)
        if not self.standardized:
            return self.net(X)
        # Otherwise:
        # Scale X:
```

```

X_scaled = self.feature_scaler.transform(X)
# Evaluate the network output - which is also scaled:
y_scaled = self.net(torch.Tensor(X_scaled))
# Scale the output back:
y = self.target_scaler.inverse_transform(y_scaled.detach().numpy())
return y

```

[ ]: # Go through the code that follows and fill in the missing parts

```

from sklearn.model_selection import train_test_split
# We need this for a progress bar:
from tqdm import tqdm

def train_net(X, y, net, reg_weight, n_batch, epochs, lr, test_size=0.33,
              standardize=True):
    """
    A function that trains a regression neural network using stochastic
    gradient descent and returns the trained network. The loss function being
    minimized is `loss_func`.

```

*Arguments:*

X	-	The observed features
y	-	The observed targets
net	-	The network you want to fit
n_batch	-	The batch size you want to use for stochastic optimization
epochs	-	How many times you want to pass over the training dataset.
lr	-	The learning rate for the stochastic optimization algorithm.
test_size	-	What % of the data should be used for testing (validation).
standardize	-	Whether or not you want to standardize the features/targets.

```

# Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33)
```

```

# Standardize the data
if standardize:
    # Build the scalers
    feature_scaler = StandardScaler().fit(X)
    target_scaler = StandardScaler().fit(y)
    # Get scaled versions of the data
    X_train_scaled = feature_scaler.transform(X_train)
    y_train_scaled = target_scaler.transform(y_train)
    X_test_scaled = feature_scaler.transform(X_test)
    y_test_scaled = target_scaler.transform(y_test)
else:
    feature_scaler = None
    target_scaler = None
    X_train_scaled = X_train
```

```

y_train_scaled = y_train
X_test_scaled = X_test
y_test_scaled = y_test

# Turn all the numpy arrays to torch tensors
X_train_scaled = torch.Tensor(X_train_scaled)
X_test_scaled = torch.Tensor(X_test_scaled)
y_train_scaled = torch.Tensor(y_train_scaled)
y_test_scaled = torch.Tensor(y_test_scaled)

# This is pytorch magic to enable shuffling of the
# training data every time we go through them
train_dataset = torch.utils.data.TensorDataset(X_train_scaled, □
y_train_scaled)
train_data_loader = torch.utils.data.DataLoader(train_dataset,
                                                batch_size=n_batch,
                                                shuffle=True)

# Create an Adam optimizing object for the neural network `net`
# with learning rate `lr`
# raise NotImplementedError('Define the optimizer object! Delete me then!')
optimizer = torch.optim.Adam(net.parameters(), lr=lr) # your code here

# This is a place to keep track of the test loss
test_loss = []

# Iterate the optimizer.
# Remember, each time we go through the entire dataset we complete an epoch
# I have wrapped the range around tqdm to give you a nice progress bar
# to look at
for e in tqdm(range(epochs)):
    # This loop goes over all the shuffled training data
    # That's why the DataLoader class of PyTorch is convenient
    for X_batch, y_batch in train_data_loader:
        # Perform a single optimization step with loss function
        # loss_func(y_batch, y_pred, reg_weight, net.parameters())
        # Hint 1: You have defined loss_func() already
        # Hint 2: Consult the hands-on activities for an example
        # Your code here
        # raise NotImplementedError('Write stochastic gradient step code!')
>Delete me then!)

        # Zero out the gradients
optimizer.zero_grad()
# Make predictions
y_pred = net(X_batch)
# Evaluate the loss

```

```

        loss = loss_func(y_batch, y_pred, reg_weight, net.parameters())
        # Backward pass
        loss.backward()
        # Make a step
        optimizer.step()

        # Evaluate the test loss and append it on the list `test_loss`
        y_pred_test = net(X_test_scaled)
        ts_loss = mse_loss(y_test_scaled, y_pred_test)
        test_loss.append(ts_loss.item())

    # Make a TrainedModel
    trained_model = TrainedModel(net, standardized=standardize,
                                 feature_scaler=feature_scaler,
                                 target_scaler=target_scaler)

    # Make sure that we return properly scaled

    # Return everything we need to analyze the results
    return trained_model, test_loss, X_train, y_train, X_test, y_test

```

Use this to test your code:

```
[ ]: # A simple one-layer network with 10 neurons
net = nn.Sequential(nn.Linear(5, 20),
                    nn.Sigmoid(),
                    nn.Linear(20, 1))

epochs = 1000
lr = 0.01
reg_weight = 0
n_batch = 100
model, test_loss, X_train, y_train, X_test, y_test = train_net(
    X,
    y,
    net,
    reg_weight,
    n_batch,
    epochs,
    lr
)
```

100% | 1000/1000 [00:31<00:00, 31.85it/s]


There are a few more things for you to do here. First, plot the evolution of the test loss as a function of the number of epochs:

```
[ ]: # Your code here
fig, ax = plt.subplots()
```

```

ax.plot(test_loss)
ax.set_xlabel('Number of Epochs')
ax.set_ylabel('Test Loss')
ax.set_title('Evolution of Test Loss')
sns.despine(trim=True)

```



Now plot the observations vs predictions plot for the training data:


```

[ ]: # Your code here
X_train_tensor = torch.Tensor(X_train)
y_train_pred = model(X_train_tensor)

# Create the observations vs. predictions plot for training data
fig, ax = plt.subplots()
ax.scatter(y_train_pred, y_train, alpha=0.7, edgecolor='k')
yys = np.linspace(y_train.min(), y_train.max(), 100)
ax.plot(yys, yys, 'r')
ax.set_xlabel('Predictions')
ax.set_ylabel('Observations')
ax.set_title('Training Data')

```


```
sns.despine(trim=True)
```



And do the observations vs predictions plot for the test data:

```
[ ]: # Your code here
X_test_tensor = torch.Tensor(X_test)
y_test_pred = model(X_test_tensor)

# Create the observations vs. predictions plot for test data
fig, ax = plt.subplots()
ax.scatter(y_test_pred, y_test, alpha=0.7, edgecolor='k')
yys = np.linspace(y_test.min(), y_test.max(), 100)
ax.plot(yys, yys, 'r')
ax.set_xlabel('Predictions')
ax.set_ylabel('Observations')
ax.set_title('Test Data')
sns.despine(trim=True)
```



**Part C.I - Investigate the effect of the batch size** For the given network, try batch sizes of 10, 25, 50, and 100 for 400 epochs. In the sample plot, show the evolution of the test loss function for each case. Which batch sizes lead to faster training times and why? Which one would you choose?

**Answer:**

```
[ ]: epochs = 400 # pick me
lr = 0.01 # pick me
reg_weight = 0 # pick me
test_losses = []
models = []
batches = [10, 25, 50, 100] # make me a list with the right batch sizes
for n_batch in batches:
    print('Training n_batch: {}'.format(n_batch))
    net = nn.Sequential(nn.Linear(5, 20),
                        nn.Sigmoid(),
                        nn.Linear(20, 1))
    model, test_loss, X_train, y_train, X_test, y_test = train_net(
        X,
```

```
    y,
    net,
    reg_weight,
    n_batch,
    epochs,
    lr
)
test_losses.append(test_loss)
models.append(model)
```


```
Training n_batch: 10
100%|     | 400/400 [01:11<00:00,  5.62it/s]

Training n_batch: 25
100%|     | 400/400 [00:32<00:00, 12.33it/s]

Training n_batch: 50
100%|     | 400/400 [00:18<00:00, 21.08it/s]

Training n_batch: 100
100%|     | 400/400 [00:12<00:00, 32.88it/s]
```

```
[ ]: fig, ax = plt.subplots(dpi=100)
for tl, n_batch in zip(test_losses, batches):
    ax.plot(tl, label='n_batch={0:d}'.format(n_batch))
ax.set_xlabel('Number of epochs')
ax.set_ylabel('Test loss')
plt.legend(loc='best');
```



*Write your observations about the batch size here*

The larger the batch size, the faster the training time. This is due to a couple of reasons. First, larger batch sizes mean there are fewer weight updates per epoch. Fewer updates means fewer forward/backward passes and fewer optimizer steps per epoch. Second, larger batches allow for better parallelism on hardware and more efficient matrix operations.

Although large batch sizes train faster, they require more epochs to converge to a minimum test loss. In contrast, smaller batches converge to lower test losses, meaning that they are more accurate.

It is clear that there is a tradeoff between training time and model accuracy. The batch size of 100 has poor accuracy compared to the other batch sizes, making it a bad choice. The batch size of 10 takes more than twice as long to train as the next slowest batch size, also making it a poor choice. The batch size of 50 reaches approximately the same minimum loss as the batch size of 25, except it is almost twice as fast to train. Therefore, I will choose a batch size of 50 because it has a good balance between speed and accuracy.

**Part C.II - Investigate the effect of the learning rate** Fix the batch size to the best one you identified in Part C.I. For the given network, try learning rates of 1, 0.1, 0.01, and 0.001 for 400 epochs. In the sample plot, show the evolution of the test loss function for each case. Does the algorithm converge for all learning rates? Which learning rate would you choose?

**Answer:**

```
[ ]: # your code here
epochs = 400
lrs = [1, 0.1, 0.01, 0.001]
reg_weight = 0
test_losses = []
models = []
batch_size = 50

for lr in lrs:
    print('Training learning rate: {}'.format(lr))
    net = nn.Sequential(nn.Linear(5, 20),
                        nn.Sigmoid(),
                        nn.Linear(20, 1))
    model, test_loss, X_train, y_train, X_test, y_test = train_net(
        X,
        y,
        net,
        reg_weight,
        batch_size,
        epochs,
        lr
    )
    test_losses.append(test_loss)
    models.append(model)
```


Training learning rate: 1  
100% | 400/400 [00:26<00:00, 14.97it/s]

Training learning rate: 0.1  
100% | 400/400 [00:19<00:00, 20.07it/s]

Training learning rate: 0.01  
100% | 400/400 [00:19<00:00, 20.04it/s]

Training learning rate: 0.001  
100% | 400/400 [00:20<00:00, 19.55it/s]

```
[ ]: fig, ax = plt.subplots(dpi=100)
for tl, lr in zip(test_losses, lrs):
    ax.plot(tl, label='Learning Rate = {}'.format(lr))
ax.set_xlabel('Number of epochs')
ax.set_ylabel('Test loss')
plt.legend(loc='best');
```



No, the algorithm does not converge for all learning rates. When the learning rate is 1, there are significant oscillations in the test loss, and the loss does not converge to a minimum. In contrast, learning rates of 0.1, 0.01, and 0.001 all steadily converge to lower losses.

The training times were approximately equal for all learning rates. When choosing the best learning rate, we must make a tradeoff between the test loss and noise. Although the learning rate of 0.001 has the least noise, it has a higher test loss compared to 0.01 and 0.1. Although the learning rate of 0.1 has the lowest test loss, it is noisier than 0.001 and 0.01. The learning rate of 0.01 has a good balance between a low test loss and minimal noise. Therefore, I will choose a learning rate of 0.01.

**Part C.III - Investigate the effect of the regularization weight** Fix the batch size to the value you selected in C.I and the learning rate to the value you selected in C.II. For the given network, try regularization weights of 0, 1e-16, 1e-12, 1e-6, and 1e-3 for 400 epochs. In the sample plot, show the evolution of the test loss function for each case. Which regularization weight seems to be the best and why?

**Answer:**

```
[ ]: # Your code here
epochs = 400
lr = 0.01
reg_weight = [0, 1e-16, 1e-12, 1e-6, 1e-3]
```

```

test_losses = []
models = []
batch_size = 50

for rw in reg_weight:
    print('Training regularization weight: {0:.0e}'.format(rw))
    net = nn.Sequential(nn.Linear(5, 20),
                        nn.Sigmoid(),
                        nn.Linear(20, 1))
    model, test_loss, X_train, y_train, X_test, y_test = train_net(
        X,
        y,
        net,
        rw,
        batch_size,
        epochs,
        lr
    )
    test_losses.append(test_loss)
    models.append(model)

```

```


Training regularization weight: 0e+00
100%| 400/400 [00:21<00:00, 18.44it/s]
Training regularization weight: 1e-16
100%| 400/400 [00:19<00:00, 20.63it/s]
Training regularization weight: 1e-12
100%| 400/400 [00:20<00:00, 19.05it/s]
Training regularization weight: 1e-06
100%| 400/400 [00:20<00:00, 19.44it/s]
Training regularization weight: 1e-03
100%| 400/400 [00:19<00:00, 20.88it/s]

```

```

[ ]: fig, ax = plt.subplots(dpi=100)
for tl, rw in zip(test_losses, reg_weight):
    ax.plot(tl, label='Regularization Weight = {0:.0e}'.format(rw))
ax.set_xlabel('Number of epochs')
ax.set_ylabel('Test loss')
plt.legend(loc='best');

```



When the regularization weight is too small, then there is not enough regularization to prevent overfitting. When the regularization weight is too large, this may lead to underfitting. Therefore, we must choose the optimal regularization weight that minimizes overfitting, underfitting, and test loss.

The regularization weight of 1e-3 has a significantly higher test loss, making it a poor choice. In contrast, the other regularization weights (0, 1e-16, 1e-12, 1e-6) all converge to low losses. To strike a balance between overfitting and underfitting, we should choose one of the middle values (1e-16 or 1e-12). For this example, I will choose 1e-16 since it has a slightly lower loss.

**Part D.I - Train a bigger network** You have developed some intuition about the parameters involved in training a network. Now, let's train a larger one. In particular, use a 5-layer deep network with 100 neurons per layer. You can use the sigmoid activation function, or you can change it to something else. Make sure you plot: - the evolution of the test loss as a function of the epochs - the observations vs predictions plot for the test data

**Answer:**

```
[ ]: # your code here
# 5-layer network with 100 neurons per layer
net = nn.Sequential(nn.Linear(5, 100),
                     nn.Sigmoid(),
```

```

        nn.Linear(100, 100),
        nn.Sigmoid(),
        nn.Linear(100, 100),
        nn.Sigmoid(),
        nn.Linear(100, 100),
        nn.Sigmoid(),
        nn.Linear(100, 1))

# Set the parameters determined from part C
epochs = 400
lr = 0.01
reg_weight = 1e-16
batch_size = 50

# Train the network
model, test_loss, X_train, y_train, X_test, y_test = train_net(
    X,
    y,
    net,
    reg_weight,
    batch_size,
    epochs,
    lr)

```


100% | 400/400 [00:38<00:00, 10.33it/s]

[ ]: *# Plot the evolution of the test loss as a function of the epochs*

```


fig, ax = plt.subplots()
ax.plot(test_loss)
ax.set_xlabel('Number of Epochs')
ax.set_ylabel('Test Loss')
ax.set_title('Evolution of Test Loss')
sns.despine(trim=True)

```



```
[ ]: # Plot the observations vs. predictions plot for test data
X_test_tensor = torch.Tensor(X_test)
y_test_pred = model(X_test_tensor)

fig, ax = plt.subplots()
ax.scatter(y_test_pred, y_test, alpha=0.7, edgecolor='k')
yys = np.linspace(y_test.min(), y_test.max(), 100)
ax.plot(yys, yys, 'r')
ax.set_xlabel('Predictions')
ax.set_ylabel('Observations')
ax.set_title('Observations vs. Predictions for Test Data')
sns.despine(trim=True)
```



**Part D.II - Make a prediction** Visualize the scaled sound level as a function of the stream velocity for a fixed frequency of 2500 Hz, a chord length of 0.1 m, a suction side displacement thickness of 0.01 m, and an angle of attack of 0, 5, and 10 degrees.

**Answer:**

This is just a check for your model. You will have to run the following code segments for the best model you have found.

```
[ ]: best_model = model # set this equal to your best model

def plot_sound_level_as_func_of_stream_vel(
    freq=2500,
    angle_of_attack=10,
    chord_length=0.1,
    suc_side_disp_thick=0.01,
    ax=None,
    label=None
):
    if ax is None:
```

```

fig, ax = plt.subplots(dpi=100)

# The velocities on which we want to evaluate the model
vel = np.linspace(X[:, 3].min(), X[:, 3].max(), 100)[:, None]

# Make the input for the model
freqs = freq * np.ones(vel.shape)
angles = angle_of_attack * np.ones(vel.shape)
chords = chord_length * np.ones(vel.shape)
sucs = suc_side_disp_thick * np.ones(vel.shape)

# Put all these into a single array
XX = np.hstack([freqs, angles, chords, vel, sucs])

ax.plot(vel, best_model(XX), label=label)


ax.set_xlabel('Velocity (m/s)')
ax.set_ylabel('Scaled sound pressure level (decibels)')

```

```

[ ]: fig, ax = plt.subplots(dpi=100)
for aofa in [0, 5, 10]:
    plot_sound_level_as_func_of_stream_vel(
        angle_of_attack=aofa,
        ax=ax,
        label='Angle of attack={0:1.2f}'.format(aofa))
plt.legend(loc='best');

```



## 1.4 Problem 2 - Classification with DNNs

Dr. Ali Lenjani kindly provided this homework problem. It is based on our joint work on this paper: [Hierarchical convolutional neural networks information fusion for activity source detection in smart buildings](#). The data come from the [Human Activity Benchmark](#) published by Dr. Juan M. Caicedo.

So the problem is as follows. You want to put sensors on a building so that it can figure out what is going on inside it. This has applications in industrial facilities (e.g., detecting if there was an accident), public infrastructure, hospitals (e.g., did a patient fall off a bed), etc. Typically, the problem is addressed using cameras. Instead of cameras, we will investigate the ability of acceleration sensors to tell us what is going on.

Four acceleration sensors have been placed in different locations in the benchmark building to record the floor vibration signals of other objects falling from several heights. A total of seven cases were considered:

- **bag-high:** 450 g bag containing plastic pieces is dropped roughly from 2.10 m
- **bag-low:** 450 g bag containing plastic pieces is dropped roughly from 1.45 m
- **ball-high:** 560 g basketball is dropped roughly from 2.10 m
- **ball-low:** 560 g basketball is dropped roughly from 1.45 m
- **j-jump:** person 1.60 m tall, 55 kg jumps approximately 12 cm high
- **d-jump:** person 1.77 m tall, 80 kg jumps approximately 12 cm high

- **w-jump:** person 1.85 m tall, 85 kg jumps approximately 12 cm high

Each of these seven cases was repeated 115 times at five different building locations. The original data are [here](#), but I have repackaged them for you in a more convenient format. Let's download them:

```
[ ]: !curl -O 'https://dl.dropboxusercontent.com/s/n8dczk7t8bx0pxi/
↳human_activity_data.npz'
```

% Total	% Received	% Xferd	Average Speed	Time	Time	Time	Current
			Dload Upload	Total	Spent	Left	Speed
100	1457	100	1457	0	0	3540	0
					--:--:--	--:--:--	--:--:--
							3545

Here is how to load the data:

```
[ ]: data = np.load('human_activity_data.npz')
```

Alternatively, manually import the data file from Google Drive:

```
[ ]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
[ ]: file_path = '/content/drive/MyDrive/Colab Notebooks/human_activity_data.npz'
data = np.load(file_path)
```


This is a Python dictionary that contains the following entries:

```
[ ]: for key in data.keys():
    print(key, ':', data[key].shape)
```

```
features : (4025, 4, 3305)
labels_1 : (4025,)
labels_2 : (4025,)
loc_ids : (4025,)
```

Let's go over these one by one. First, the **features**. These are the acceleration sensor measurements. Here is how you visualize them:

```
[ ]: fig, ax = plt.subplots(4, 1, dpi=100)
# Loop over sensors
for j in range(4):
    ax[j].plot(data['features'][0, j])
ax[-1].set_xlabel('Timestep')
ax[-1].set_ylabel('Acceleration');
```





The second key, `labels_1`, is a bunch of integers ranging from 0 to 2 indicating whether the entry corresponds to a “bag,” a “ball” or a “jump.” For your reference, the correspondence is:


```
[ ]: LABELS_1_TO_TEXT = {
    0: 'bag',
    1: 'ball',
    2: 'jump'
}
```


And here are a few examples:


```
[ ]: for _ in range(5):
    i = np.random.randint(0, data['features'].shape[0])
    fig, ax = plt.subplots(4, 1, dpi=100)
    for j in range(4):
        ax[j].plot(data['features'][i, j])
        ax[-1].set_xlabel('Timestep')
        ax[-1].set_ylabel('Acceleration')
        ax[0].set_title('Label: {}:{} -> {}'.format(data['labels_1'][i],
                                                       □
                                                       ↵LABELS_1_TO_TEXT[data['labels_1'][i]])))
```











The array `labels_2` includes integers from 0 to 6 indicating the detailed label of the experiment. The correspondence between integers and text labels is:

```
[ ]: LABELS_2_TO_TEXT = {
    0: 'bag-high',
    1: 'bag-low',
    2: 'ball-high',
    3: 'ball-low',
    4: 'd-jump',
    5: 'j-jump',
    6: 'w-jump'
}
```

Finally, the field `loc_ids` takes values from 0 to 4 indicating five distinct locations in the building. Before moving forward with the questions, let's extract the data in a more convenient form:

```
[ ]: # The features
X = data['features']
# The labels_1
y1 = data['labels_1']
```

```

# The labels_2
y2 = data['labels_2']
# The locations
y3 = data['loc_ids']

```

#### 1.4.1 Part A - Train a CNN to predict the high-level type of observation (bag, ball, or jump)

Fill in the blanks in the code blocks below to train a classification neural network that will take you from the four acceleration sensor data to the high-level type of each observation. You can keep the network structure fixed, but you can experiment with the learning rate, the number of epochs, or anything else. Just keep in mind that for this particular dataset, it is possible to hit an accuracy of almost 100%.

**Answer:**

The first thing that we need to do is pick a neural network structure. Let's use 1D convolutional layers at the very beginning. These are the same as the 2D (image) convolutional layers but in 1D. The reason I am proposing this is that the convolutional layers are invariant to small translations of the acceleration signal (just like the labels are). Here is what I propose:

```

[ ]: import torch
import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):
    def __init__(self, num_labels=3):
        super(Net, self).__init__()
        # A convolutional layer:
        # 3 = input channels (sensors),
        # 6 = output channels (features),
        # 5 = kernel size
        self.conv1 = nn.Conv1d(4, 8, 10)
        # A 2 x 2 max pooling layer - we are going to use it two times
        self.pool = nn.MaxPool1d(5)
        # Another convolutional layer
        self.conv2 = nn.Conv1d(8, 16, 5)
        # Some linear layers
        self.fc1 = nn.Linear(16 * 131, 200)
        self.fc2 = nn.Linear(200, 50)
        self.fc3 = nn.Linear(50, num_labels)

    def forward(self, x):
        # This function implements your network output
        # Convolutional layer, followed by relu, followed by max pooling
        x = self.pool(F.relu(self.conv1(x)))
        # Same thing
        x = self.pool(F.relu(self.conv2(x)))

```

```

# Flattening the output of the convolutional layers
x = x.view(-1, 16 * 131)
# Go through the first dense linear layer followed by relu
x = F.relu(self.fc1(x))
# Through the second dense layer
x = F.relu(self.fc2(x))
# Finish up with a linear transformation
x = self.fc3(x)
return x

```

[ ]: # You can make the network like this:

```
net = Net(3)
```

Now, you need to pick the right loss function for classification tasks:

[ ]: # your code here

```
cnn_loss_func = nn.CrossEntropyLoss()
```

Just like before, let's organize our training code in a convenient function that allows us to play with the parameters of training. Fill in the missing code.

```

[ ]: def train_cnn(X, y, net, n_batch, epochs, lr, test_size=0.33):
    """
    A function that trains a regression neural network using stochastic
    gradient descent and returns the trained network. The loss function being
    minimized is `loss_func`.
    """

    Parameters:

    X      -   The observed features
    y      -   The observed targets
    net    -   The network you want to fit
    n_batch -   The batch size you want to use for stochastic optimization
    epochs  -   How many times you want to pass over the training dataset.
    lr      -   The learning rate for the stochastic optimization algorithm.
    test_size -   What percent of data should be used for testing (validation).
    """

    # Split the data
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33)

    # Turn all the numpy arrays to torch tensors
    X_train = torch.Tensor(X_train)
    X_test = torch.Tensor(X_test)
    y_train = torch.LongTensor(y_train)
    y_test = torch.LongTensor(y_test)

    # This is pytorch magick to enable shuffling of the
    # training data every time we go through them

```

```

train_dataset = torch.utils.data.TensorDataset(X_train, y_train)
train_data_loader = torch.utils.data.DataLoader(train_dataset,
                                              batch_size=n_batch,
                                              shuffle=True)

# Create an Adam optimizing object for the neural network `net`
# with learning rate `lr`
# raise NotImplementedError('Define the optimizer object! Delete me then!')
optimizer = torch.optim.Adam(net.parameters(), lr=lr) # your code here

# This is a place to keep track of the test loss
test_loss = []
# This is a place to keep track of the accuracy on each epoch
accuracy = []

# Iterate the optimizer.
# Remember, each time we go through entire dataset we complete an `epoch`
# I have wrapped the range around tqdm to give you a nice progress bar
# to look at
for e in range(epochs):
    # This loop goes over all the shuffled training data
    # That's why the DataLoader class of PyTorch is convenient
    for X_batch, y_batch in train_data_loader:
        # Perform a single optimization step with loss function
        # cnn_loss_func(y_batch, y_pred, reg_weight)
        # Hint 1: You have defined cnn_loss_func() already
        # Hint 2: Consult the hands-on activities for an example
        # your code here
        # raise NotImplementedError('Write stochastic gradient step code! ↴Delete me then!')

        optimizer.zero_grad()
        y_pred = net(X_batch)
        loss = cnn_loss_func(y_pred, y_batch)
        loss.backward()
        optimizer.step()

    # Evaluate the test loss and append it on the list `test_loss`
    y_pred_test = net(X_test)
    ts_loss = cnn_loss_func(y_pred_test, y_test)
    test_loss.append(ts_loss.item())
    # Evaluate the accuracy
    _, predicted = torch.max(y_pred_test.data, 1)
    correct = (predicted == y_test).sum().item()
    accuracy.append(correct / y_test.shape[0])
    # Print something about the accuracy
    print('Epoch {0:d}: accuracy = {1:1.5f}%'.format(e+1, accuracy[-1]))

```

```

trained_model = net

# Return everything we need to analyze the results
return trained_model, test_loss, accuracy, X_train, y_train, X_test, y_test

```

Now experiment with the epochs, the learning rate, and the batch size until this works.


```
[ ]: epochs = 20
lr = 0.001
n_batch = 30
trained_model, test_loss, accuracy, X_train, y_train, X_test, y_test =  

    ↪train_cnn(X, y1, net, n_batch, epochs, lr)
```

Epoch 1: accuracy = 0.43115%  
 Epoch 2: accuracy = 0.70956%  
 Epoch 3: accuracy = 0.96614%  
 Epoch 4: accuracy = 0.97065%  
 Epoch 5: accuracy = 0.98947%  
 Epoch 6: accuracy = 0.99022%  
 Epoch 7: accuracy = 0.99248%  
 Epoch 8: accuracy = 0.99473%  
 Epoch 9: accuracy = 0.99172%  
 Epoch 10: accuracy = 0.99323%  
 Epoch 11: accuracy = 0.99624%  
 Epoch 12: accuracy = 0.99774%  
 Epoch 13: accuracy = 0.99699%  
 Epoch 14: accuracy = 0.99699%  
 Epoch 15: accuracy = 0.99549%  
 Epoch 16: accuracy = 0.99850%  
 Epoch 17: accuracy = 0.99699%  
 Epoch 18: accuracy = 0.99699%  
 Epoch 19: accuracy = 0.99850%  
 Epoch 20: accuracy = 0.99925%


Plot the evolution of the test loss as a function of epochs.

```
[ ]: fig, ax = plt.subplots(dpi=100)
ax.plot(test_loss)
ax.set_xlabel('Number of epochs')
ax.set_ylabel('Test loss');
ax.set_xticks(np.arange(0, epochs + 1, 5));
```



Plot the evolution of the accuracy as a function of epochs.

```
[ ]: fig, ax = plt.subplots(dpi=100)
ax.plot(accuracy)
ax.set_xlabel('Number of epochs')
ax.set_ylabel('Accuracy');
ax.set_xticks(np.arange(0, epochs + 1, 5));
```




Plot the confusion matrix.

```
[ ]: from sklearn.metrics import confusion_matrix
# Predict on the test data
y_pred_test = trained_model(X_test)
# Remember that the prediction is probabilistic
# We need to simply pick the label with the highest probability:
_, y_pred_labels = torch.max(y_pred_test, 1)
# Here is the confusion matrix:
cf_matrix = confusion_matrix(y_test, y_pred_labels)
```

```
[ ]: sns.heatmap(cf_matrix/np.sum(cf_matrix), annot=True,
                fmt='.2%', cmap='Blues',
                xticklabels=LABELS_1_TO_TEXT.values(),
                yticklabels=LABELS_1_TO_TEXT.values());
```



#### 1.4.2 Part B - Train a CNN to predict the the low-level type of observation (bag-high, bag-low, etc.)

Repeat what you did above for  $y_2$ .

**Answer:**

```
[ ]: # your code here
# Make the network
net = Net(7)

# Loss function
cnn_loss_func = nn.CrossEntropyLoss()

# Set the parameters and train the model
epochs = 100
lr = 0.001
n_batch = 30
trained_model, test_loss, accuracy, X_train, y_train, X_test, y_test = ↵
    train_cnn(X, y2, net, n_batch, epochs, lr)
```

Epoch 1: accuracy = 0.14296%

Epoch 2: accuracy = 0.19112%


Epoch 3: accuracy = 0.43491%  
Epoch 4: accuracy = 0.58315%  
Epoch 5: accuracy = 0.58239%  
Epoch 6: accuracy = 0.67645%  
Epoch 7: accuracy = 0.69375%  
Epoch 8: accuracy = 0.69375%  
Epoch 9: accuracy = 0.73589%  
Epoch 10: accuracy = 0.72310%  
Epoch 11: accuracy = 0.73965%  
Epoch 12: accuracy = 0.72084%  
Epoch 13: accuracy = 0.78029%  
Epoch 14: accuracy = 0.74944%  
Epoch 15: accuracy = 0.77502%  
Epoch 16: accuracy = 0.75771%  
Epoch 17: accuracy = 0.78555%  
Epoch 18: accuracy = 0.81189%  
Epoch 19: accuracy = 0.84500%  
Epoch 20: accuracy = 0.84349%  
Epoch 21: accuracy = 0.79834%  
Epoch 22: accuracy = 0.83446%  
Epoch 23: accuracy = 0.84876%  
Epoch 24: accuracy = 0.87284%  
Epoch 25: accuracy = 0.88488%  
Epoch 26: accuracy = 0.88412%  
Epoch 27: accuracy = 0.88488%  
Epoch 28: accuracy = 0.89165%  
Epoch 29: accuracy = 0.89090%  
Epoch 30: accuracy = 0.89165%  
Epoch 31: accuracy = 0.89691%  
Epoch 32: accuracy = 0.89842%  
Epoch 33: accuracy = 0.89767%  
Epoch 34: accuracy = 0.90218%  
Epoch 35: accuracy = 0.91046%  
Epoch 36: accuracy = 0.90444%  
Epoch 37: accuracy = 0.90369%  
Epoch 38: accuracy = 0.86983%  
Epoch 39: accuracy = 0.90745%  
Epoch 40: accuracy = 0.89917%  
Epoch 41: accuracy = 0.88713%  
Epoch 42: accuracy = 0.90444%  
Epoch 43: accuracy = 0.89992%  
Epoch 44: accuracy = 0.89992%  
Epoch 45: accuracy = 0.89466%  
Epoch 46: accuracy = 0.88638%  
Epoch 47: accuracy = 0.89240%  
Epoch 48: accuracy = 0.90971%  
Epoch 49: accuracy = 0.90820%  
Epoch 50: accuracy = 0.92250%

Epoch 51: accuracy = 0.90293%  
Epoch 52: accuracy = 0.91272%  
Epoch 53: accuracy = 0.91949%  
Epoch 54: accuracy = 0.90670%  
Epoch 55: accuracy = 0.90971%  
Epoch 56: accuracy = 0.89466%  
Epoch 57: accuracy = 0.91196%  
Epoch 58: accuracy = 0.86456%  
Epoch 59: accuracy = 0.92024%  
Epoch 60: accuracy = 0.92476%  
Epoch 61: accuracy = 0.87509%  
Epoch 62: accuracy = 0.91196%  
Epoch 63: accuracy = 0.90519%  
Epoch 64: accuracy = 0.92626%  
Epoch 65: accuracy = 0.92250%  
Epoch 66: accuracy = 0.89315%  
Epoch 67: accuracy = 0.91648%  
Epoch 68: accuracy = 0.90369%  
Epoch 69: accuracy = 0.88412%  
Epoch 70: accuracy = 0.90218%  
Epoch 71: accuracy = 0.88412%  
Epoch 72: accuracy = 0.91648%  
Epoch 73: accuracy = 0.91272%  
Epoch 74: accuracy = 0.91046%  
Epoch 75: accuracy = 0.92927%  
Epoch 76: accuracy = 0.92325%  
Epoch 77: accuracy = 0.90670%  
Epoch 78: accuracy = 0.91949%  
Epoch 79: accuracy = 0.92551%  
Epoch 80: accuracy = 0.92099%  
Epoch 81: accuracy = 0.91121%  
Epoch 82: accuracy = 0.89691%  
Epoch 83: accuracy = 0.91347%  
Epoch 84: accuracy = 0.88939%  
Epoch 85: accuracy = 0.92476%  
Epoch 86: accuracy = 0.91497%  
Epoch 87: accuracy = 0.89541%  
Epoch 88: accuracy = 0.90895%  
Epoch 89: accuracy = 0.92626%  
Epoch 90: accuracy = 0.91949%  
Epoch 91: accuracy = 0.89541%  
Epoch 92: accuracy = 0.91272%  
Epoch 93: accuracy = 0.90820%  
Epoch 94: accuracy = 0.92927%  
Epoch 95: accuracy = 0.93153%  
Epoch 96: accuracy = 0.92852%  
Epoch 97: accuracy = 0.92927%  
Epoch 98: accuracy = 0.93153%

```
Epoch 99: accuracy = 0.92852%
Epoch 100: accuracy = 0.92325%
```


Plot the evolution of the test loss as a function of epochs:

```
[ ]: # Plot the evolution of the test loss as a function of epochs
fig, ax = plt.subplots(dpi=100)
ax.plot(test_loss)
ax.set_xlabel('Number of epochs')
ax.set_ylabel('Test loss');
ax.set_xticks(np.arange(0, epochs + 1, 20));
```



Plot the evolution of the accuracy as a function of epochs:

```
[ ]: # Plot the evolution of the accuracy as a function of epochs
fig, ax = plt.subplots(dpi=100)
ax.plot(accuracy)
ax.set_xlabel('Number of epochs')
ax.set_ylabel('Accuracy');
ax.set_xticks(np.arange(0, epochs + 1, 20));
```



Plot the confusion matrix:

```
[ ]: # Plot the confusion matrix
# Predict on the test data
y_pred_test = trained_model(X_test)
# We need to pick the label with the highest probability
_, y_pred_labels = torch.max(y_pred_test, 1)
# Here is the confusion matrix:
cf_matrix = confusion_matrix(y_test, y_pred_labels)
# Create the plot
sns.heatmap(cf_matrix/np.sum(cf_matrix), annot=True,
            fmt='%.2%', cmap='Blues',
            xticklabels=LABELS_2_TO_TEXT.values(),
            yticklabels=LABELS_2_TO_TEXT.values());
```

